

Banking Core - Database Normalization Analysis

Schema Design Justification & Normal Form Compliance

Version: 1.0.0

Last Updated: January 2026

Target Audience: DBMS Undergraduate Students, Faculty

Database: MySQL 8.0 (InnoDB)

1. Introduction to Normalization in Banking Systems

1.1 Purpose of Normalization

Database normalization is the process of organizing a relational database to:

- 1. Minimize redundancy** – Store each fact only once
- 2. Prevent anomalies** – Avoid insert, update, and delete problems
- 3. Ensure data integrity** – Maintain consistency across tables

1.2 Banking-Specific Considerations

Financial databases have unique requirements that influence normalization decisions:

Requirement	Impact on Design
Transaction Integrity	Strict normalization for financial records
Audit Compliance	Immutable audit trails (append-only)
Performance	Strategic denormalization for balance lookups
Regulatory Requirements	Complete history preservation

1.3 Normal Forms Overview

Normal Form	Rule
1NF	Atomic values, no repeating groups
2NF	1NF + No partial dependencies on composite keys
3NF	2NF + No transitive dependencies
BCNF	Every determinant is a candidate key

2. Functional Dependencies Analysis

A **functional dependency** $X \rightarrow Y$ means that for each value of X, there is exactly one value of Y.

2.1 Core Entity Tables

roles

```
id → {code, name, description, permissions, is_system, created_at, updated_at}  
code → {id, name, description, permissions, is_system}
```

Candidate Keys: {id}, {code}

users

```
id → {email, password_hash, first_name, last_name, role_id, status, last_login_at, ...}  
email → {id, password_hash, first_name, last_name, role_id, ...}
```

Candidate Keys: {id}, {email}

customers

```
id → {customer_number, email, first_name, last_name, phone, national_id, ...}  
customer_number → {id, email, first_name, ...}  
email → {id, customer_number, first_name, ...}
```

Candidate Keys: {id}, {customer_number}, {email}

2.2 Account Tables

account_types

```
id → {code, name, description, min_balance, is_active}  
code → {id, name, description, min_balance, is_active}
```

Candidate Keys: {id}, {code}

accounts

```
id → {account_number, customer_id, account_type_id, status, opened_at, ...}  
account_number → {id, customer_id, account_type_id, status, ...}
```

Candidate Keys: {id}, {account_number}

account_balances

```
account_id → {available_balance, pending_balance, hold_balance, currency, version, ...}
```

Candidate Keys: {account_id}

2.3 Transaction Tables

transaction_types

```
id → {code, name, description, requires_approval}  
code → {id, name, description, requires_approval}
```

Candidate Keys: {id}, {code}

transactions

```
id → {transaction_reference, transaction_type_id, amount, currency, status,  
      source_account_id, destination_account_id, processed_at, ...}  
transaction_reference → {id, transaction_type_id, amount, ...}
```

Candidate Keys: {id}, {transaction_reference}

ledger_entries

```
id → {transaction_id, account_id, entry_type, amount, currency, balance_after,  
description, entry_date, created_at}
```

Candidate Keys: {id}

2.4 Audit Tables

transaction_audit

```
id → {ledger_entry_id, transaction_id, account_id, entry_type, amount,  
balance_after, audit_timestamp}  
ledger_entry_id → {id, transaction_id, account_id, ...}
```

Candidate Keys: {id}, {ledger_entry_id}

audit_logs

```
id → {actor_id, actor_type, actor_role, action_type, entity_type, entity_id,  
before_state, after_state, metadata, created_at}
```

Candidate Keys: {id}

3. First Normal Form (1NF) Compliance

3.1 Requirements for 1NF

1. Each column contains only **atomic (indivisible) values**
2. Each column contains values of a **single type**
3. Each row is **uniquely identifiable** (has a primary key)
4. No **repeating groups** or arrays within a column

3.2 Table-by-Table 1NF Analysis

Table	Atomic Values	Single Type	PK Exists	No Repeating Groups	1NF?
roles	✓	✓	✓ (id)	⚠️ See Note 1	✓
users	✓	✓	✓ (id)	✓	✓
customers	✓	✓	✓ (id)	✓	✓
account_types	✓	✓	✓ (id)	✓	✓
accounts	✓	✓	✓ (id)	✓	✓
account_balances	✓	✓	✓ (account_id)	✓	✓
transaction_types	✓	✓	✓ (id)	✓	✓
transactions	✓	✓	✓ (id)	✓	✓
ledger_entries	✓	✓	✓ (id)	✓	✓
transaction_audit	✓	✓	✓ (id)	✓	✓
audit_logs	✓	✓	✓ (id)	⚠️ See Note 2	✓
system_config	✓	✓	✓ (id)	✓	✓
idempotency_keys	✓	✓	✓ (key)	⚠️ See Note 3	✓
events	✓	✓	✓ (id)	⚠️ See Note 3	✓
outbox	✓	✓	✓ (id)	⚠️ See Note 3	✓

3.3 Notes on JSON Columns

Note 1 (roles.permissions): Contains JSON array like `["read:*", "write:"]`

- **Justification:** While JSON could be considered non-atomic, MySQL 8 treats JSON as a first-class type. Permissions are read as a whole set, not queried individually.
- **Alternative:** A separate `role_permissions` junction table would add complexity for minimal benefit.
- **Conclusion:** Acceptable trade-off; still considered 1NF compliant.

Note 2 (audit_logs.before_state, after_state, metadata): Contains JSON objects

- **Justification:** Audit logs capture arbitrary entity states. A fully normalized design would require a separate audit table per entity type.
- **Conclusion:** JSON provides flexibility for audit trails; 1NF compliant.

Note 3 (events/outbox.payload): Contains JSON event data

- **Justification:** Event payloads vary by event type. Normalizing would require dozens of event-specific tables.
- **Conclusion:** Standard practice for event sourcing; 1NF compliant.

4. Second Normal Form (2NF) Compliance

4.1 Requirements for 2NF

1. Must be in **1NF**
2. **No partial dependencies** – All non-key attributes must depend on the entire primary key

4.2 Analysis of Composite Keys

2NF is primarily concerned with tables that have **composite primary keys**. In Banking Core:

Table	Primary Key	Type
All core tables	Single column (id)	Simple
account_balances	account_id	Simple
idempotency_keys	idempotency_key	Simple

Finding: All tables use **single-column primary keys**, so **partial dependencies are impossible**.

4.3 2NF Compliance Summary

Table	1NF	Single-Column PK	Partial Dependencies	2NF?
roles	✓	✓	N/A	✓
users	✓	✓	N/A	✓

Table	1NF	Single-Column PK	Partial Dependencies	2NF?
customers	✓	✓	N/A	✓
account_types	✓	✓	N/A	✓
accounts	✓	✓	N/A	✓
account_balances	✓	✓	N/A	✓
transaction_types	✓	✓	N/A	✓
transactions	✓	✓	N/A	✓
ledger_entries	✓	✓	N/A	✓
transaction_audit	✓	✓	N/A	✓
audit_logs	✓	✓	N/A	✓
system_config	✓	✓	N/A	✓
idempotency_keys	✓	✓	N/A	✓
events	✓	✓	N/A	✓
outbox	✓	✓	N/A	✓

Result: All tables satisfy **2NF** by design (single-column PKs).

5. Third Normal Form (3NF) Compliance

5.1 Requirements for 3NF

1. Must be in **2NF**
2. **No transitive dependencies** – Non-key attributes must not depend on other non-key attributes

5.2 Transitive Dependency Analysis

A transitive dependency exists when: $A \rightarrow B \rightarrow C$ (A determines B, B determines C, but B is not a candidate key)

users table

`id → role_id → role_name (via JOIN)`

- **Analysis:** `role_name` is NOT stored in users table
- **Stored:** Only `role_id` (FK reference)
- **Result:** No redundancy, 3NF compliant

accounts table

`id → customer_id → customer_name (via JOIN)`

`id → account_type_id → account_type_name (via JOIN)`

- **Analysis:** Customer and type names are NOT stored in accounts
- **Stored:** Only foreign key IDs
- **Result:** 3NF compliant

ledger_entries table

`id → transaction_id → transaction_amount (via JOIN)`

`id → account_id → account_number (via JOIN)`

- **Analysis:** Redundant data NOT stored
- **Stored:** Only references
- **Result:** 3NF compliant

5.3 Controlled Denormalization (3NF Exceptions)

Two tables contain **intentional denormalization** for performance:

account_balances.available_balance

Normalized approach:
`SUM(ledger_entries.amount) WHERE entry_type = 'CREDIT'`
`- SUM(ledger_entries.amount) WHERE entry_type = 'DEBIT'`

Denormalized: `account_balances.available_balance (cached value)`

Justification:

- Balance queries happen on every transaction

- Summing 100,000+ ledger entries per query is O(n) expensive
- Cached balance provides O(1) lookup
- **Integrity maintained:** Stored procedures update both ledger AND balance atomically

ledger_entries.balance_after

Normalized approach: SUM all prior entries to calculate balance at any point

Denormalized: balance_after stored directly in each entry

Justification:

- Historical balance queries are common (statements, audits)
- Recalculating from genesis would be extremely expensive
- balance_after is calculated and stored at insert time
- **Never updated** – append-only table ensures integrity

5.4 3NF Compliance Summary

Table	2NF	Transitive Dependencies	Exception Justified	3NF?
roles	✓	None	N/A	✓
users	✓	None	N/A	✓
customers	✓	None	N/A	✓
account_types	✓	None	N/A	✓
accounts	✓	None	N/A	✓
account_balances	✓	Cached balance	✓ Performance	⚠ Denormalized
transaction_types	✓	None	N/A	✓
transactions	✓	None	N/A	✓
ledger_entries	✓	balance_after	✓ Historical queries	⚠ Denormalized
transaction_audit	✓	None	N/A	✓

Table	2NF	Transitive Dependencies	Exception Justified	3NF?
audit_logs	✓	None	N/A	✓
system_config	✓	None	N/A	✓
idempotency_keys	✓	None	N/A	✓
events	✓	None	N/A	✓
outbox	✓	None	N/A	✓

Result: 13/15 tables fully 3NF compliant. 2 tables have **justified denormalization**.

6. BCNF (Boyce-Codd Normal Form) Consideration

6.1 Requirements for BCNF

- For every functional dependency $X \rightarrow Y$, X must be a **superkey**
- Stricter than 3NF: addresses anomalies when multiple candidate keys overlap

6.2 BCNF Analysis

Table	Candidate Keys	All Determinants Superkeys?	BCNF?
roles	{id}, {code}	✓	✓
users	{id}, {email}	✓	✓
customers	{id}, {customer_number}, {email}	✓	✓
account_types	{id}, {code}	✓	✓
accounts	{id}, {account_number}	✓	✓
account_balances	{account_id}	✓	✓
transaction_types	{id}, {code}	✓	✓

Table	Candidate Keys	All Determinants Superkeys?	BCNF?
transactions	{id}, {transaction_reference}	✓	✓
ledger_entries	{id}	✓	✓
transaction_audit	{id}, {ledger_entry_id}	✓	✓
audit_logs	{id}	✓	✓
system_config	{id}, {config_key}	✓	✓
idempotency_keys	{idempotency_key}	✓	✓
events	{id}	✓	✓
outbox	{id}	✓	✓

Result: All tables satisfy **BCNF**.

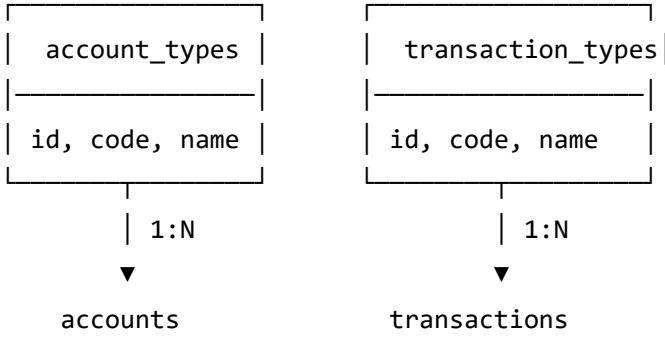
7. Schema Design Justification

7.1 Separation of Concerns

Design Decision	Justification
users vs customers	Staff and customers have different attributes, workflows, and security models
transactions vs ledger_entries	Transaction is a business event; ledger entries are accounting records
accounts vs account_balances	Account metadata (static) separated from balance (dynamic)
transaction_audit vs ledger_entries	Audit copy provides regulatory compliance verification

7.2 Reference Table Pattern

Lookup/reference tables are normalized separately:



Benefits:

- Type names updated in one place
- New types added without schema changes
- Consistent categorization across system

7.3 Double-Entry Ledger Design

transactions (1) ———< ledger_entries (N)

Each transaction generates exactly 2 ledger entries:

- **DEBIT** entry (money leaving account)
- **CREDIT** entry (money entering account)

Normalization Benefit: Transaction metadata stored once, not duplicated per entry.

8. Anomaly Prevention

8.1 Insert Anomalies

Scenario	How Schema Prevents
Adding account without customer	customer_id NOT NULL enforces relationship
Adding transaction without type	transaction_type_id NOT NULL required
Adding ledger entry without transaction	transaction_id NOT NULL enforces parent

8.2 Update Anomalies

Scenario	How Schema Prevents
Renaming account type	Update once in <code>account_types</code> , all accounts inherit via FK
Changing role permissions	Update once in <code>roles</code> , all users inherit
Customer address change	Single record in <code>customers</code> , no duplication

8.3 Delete Anomalies

Scenario	How Schema Prevents
Deleting customer with accounts	Application-level FK check prevents orphans
Deleting transaction type in use	Referenced transactions would be orphaned (blocked)
Ledger entries never deleted	Append-only design preserves history

9. Trade-offs & Performance Considerations

9.1 Normalization vs. Performance Trade-offs

Trade-off	Decision	Reason
Balance calculation	Denormalized (cached)	O(1) vs O(n) lookup
Transaction type name	Normalized (FK join)	Infrequent updates, small table
Customer in ledger	Normalized (via account FK)	Avoids massive redundancy
Historical balance	Denormalized (<code>balance_after</code>)	Statement generation performance

9.2 Join Performance Analysis

Query Pattern	Tables Joined	Index Support	Performance
Get account balance	<code>accounts</code> + <code>account_balances</code>	PK lookup	O(1)

Query Pattern	Tables Joined	Index Support	Performance
List customer accounts	accounts WHERE customer_id	FK index	O(log n)
Transaction history	transactions + ledger_entries	FK index	O(log n)
Audit trail	audit_logs	Multiple indexes	O(log n)

9.3 Index Strategy

Table	Index	Purpose
users	email (unique)	Login lookup
accounts	customer_id	List customer's accounts
accounts	account_number (unique)	Transfer destination lookup
ledger_entries	transaction_id	Transaction details
ledger_entries	account_id, entry_date	Statement queries
audit_logs	(actor_id, actor_type)	Who did what
audit_logs	(entity_type, entity_id)	What happened to entity
audit_logs	created_at	Time-based queries

10. Conclusion

10.1 Normalization Summary

Normal Form	Tables Fully Compliant	Tables with Justified Exceptions
1NF	15/15 (100%)	0
2NF	15/15 (100%)	0
3NF	13/15 (87%)	2 (account_balances, ledger_entries)
BCNF	15/15 (100%)	0

10.2 Key Design Strengths

1. **Clean Separation** – Identity, accounts, transactions, and audit cleanly separated
2. **Minimal Redundancy** – Data stored once, referenced via FKs
3. **Controlled Denormalization** – Only where performance requires it
4. **Audit Integrity** – Immutable audit trails for regulatory compliance
5. **Flexible Extensions** – JSON columns for variable-structure data

10.3 Anomaly Prevention Summary

Anomaly Type	Status
Insert Anomalies	<input checked="" type="checkbox"/> Prevented via NOT NULL constraints
Update Anomalies	<input checked="" type="checkbox"/> Prevented via normalized reference tables
Delete Anomalies	<input checked="" type="checkbox"/> Prevented via FK enforcement + append-only tables

10.4 Final Assessment

The Banking Core database schema demonstrates **strong normalization discipline** appropriate for a financial system:

- Core transactional tables are in **3NF/BCNF**
- Denormalization is **intentional and justified** for performance-critical paths
- **Audit compliance** is maintained through append-only ledger and separate audit tables
- The design **prevents all three types of anomalies** while supporting high-performance banking operations

This document is suitable for DBMS coursework, viva examinations, and academic project reports.